

CHANCE PILOT #2

Does the AI Model Change the Strategic Decision?

A controlled comparison of GPT-5.6 Sol, Claude Sonnet 5, Gemini 3.6 Flash, and Codex GPT-5.6 Sol Medium

Same product evidence. Same strategic question. Different decision-makers.

18 August 2026

Executive Summary

Question. When AI is used for product strategy, which judgments remain stable across models, and which depend on the model making the decision?

Result. The four models converged more strongly on what Chance should not rely on than on what it should become. All four moved away from treating generic single-image recognition or explanation as the strongest commercial asset. They then diverged substantially on the hidden asset, target market, buyer, and commercialization path.

The models were more consistent in rejecting Chance's apparent core than in identifying what should replace it.

The divergence began before the final recommendation. The models assigned commercial importance to different underlying assets: GPT emphasized longitudinal visual workflow, Claude persistent memory and personalization, Gemini visual observation and vocabulary guidance, and Codex repeatable photo-to-output workflows. Those different asset identifications then propagated into different users, payers, recurring-value mechanisms, and business models.

This pilot does not identify a winning model or a correct strategy. It shows where strategic judgment was robust to model choice and where it was model-sensitive under a fixed evidence base and a fixed question.

1. Experimental Design

Pilot #2 builds on one earlier Chance experiment and a follow-up analysis of that experiment. The sequence is:

- Pilot #1: changing the layer of evidence available to AI changed the product it appeared to see and the opportunities it identified.
- Follow-up analysis: AI could identify evidence that a capability was weak yet still preserve that capability in its strategic recommendation.
- Pilot #2: evidence and strategic question were held constant while the model making the strategic judgment changed.

Models compared: GPT-5.6 Sol; Claude Sonnet 5; Gemini 3.6 Flash; Codex using GPT-5.6 Sol Medium.

Controlled inputs

GPT-5.6 Sol, Claude Sonnet 5, and Gemini 3.6 Flash received the same six-page Neutral Evidence Packet in fresh contexts. The packet compressed three role-played first-use sessions and a persona-free capability audit while excluding prior strategy recommendations, preferred buyers, opportunity rankings, and business-model conclusions.

The same strategy prompt was then used. External search and additional Chance investigation were prohibited, and the models were instructed to distinguish direct evidence from inference. Each model ran without seeing the other models' answers.

Codex (GPT-5.6 Sol Medium) was handled differently because it had already operated Chance directly during the earlier capability-discovery phase and had already produced a strategy answer. That pre-existing answer was added only after the GPT-5.6 Sol, Claude Sonnet 5, and Gemini 3.6 Flash runs. It was not shown to them.

What was compared

The comparison focused on judgment layers, not prose style or model “intelligence”: strongest capability or bundle; top commercial opportunities; user and payer; B2C vs. B2B/B2B2C direction; recurring value; differentiation from generic multimodal AI; keep/kill decisions; underrated and overrated features; non-obvious opportunity; failure risk; and the evidence used to justify those decisions.

2. Evidence Base in Brief

The common packet contained enough evidence to make several plausible strategic readings possible without embedding a recommended strategy. The most decision-relevant observations were:

- P01 understood the product and found some output interesting, but identified no realistic reason to reopen within seven days and expected simple identification to remain with Google Lens and deeper questioning with ChatGPT or Gemini.
- P02 valued the conversion of vague visual impressions into vocabulary and points of attention, but a serious artwork misidentification caused a sharp trust drop. Chance built detailed interpretation on the wrong artist and title.
- P03 found OOTD understandable but generic; footwear was misclassified and the selected “Campus” context was weakly reflected.

- The audit found a broader product layer: automatic Snap history, cross-image statements, Patterns/Taste Profile, Weekly Recap, AI Memory Import, personalization, agent discovery, visual/card generation, and sharing.
- Reliability problems also appeared in shopping/price matching, target selection, and source quality. A known R\$280 cushion produced suggestions around R\$19.90–42.00.

The full Neutral Evidence Packet is reproduced in Appendix A.

3. Four Models, Four Commercial Directions

The strongest opportunity selected by each model was materially different.

Model	Top direction	Core asset identified	Primary payer	Orientation
GPT-5.6 Sol	Visual Research Workspace for design / merchandising teams	Longitudinal visual workflow: capture → archive → cross-image analysis → reusable output	Brand / design / retail teams	B2B
Claude Sonnet 5	Persistent Visual Life-Journal	Persistent memory, cross-Snap history, personalization	Consumer subscriber	B2C
Gemini 3.6 Flash	Art / Design visual-learning agent	Visual observation, vocabulary, guided interpretation	Consumer and/or education institution	B2C / B2B2C
Codex (GPT-5.6 Sol Medium)	Brand visual campaign platform	Repeatable photo-to-output workflow with agents, branded artifacts, sharing and history	Brands	B2B / B2B2C

The disagreement was not merely about ranking. GPT and Codex both moved toward professional workflows, but they selected different jobs. Claude retained a consumer model. Gemini elevated an education-oriented interpretation that the others did not place first.

Top-five opportunity patterns

Across the full ranked lists, recurring clusters appeared, but no single commercialization path dominated all four models:

- Longitudinal memory / Journal / cross-Snap use appeared repeatedly in GPT, Claude, and Codex, and remained part of Gemini's retained bundle.
- Professional visual workflows appeared in GPT and Codex, but with different task definitions: research/merchandising vs. campaigns, field operations, resale intake, clienteling, and beauty-service records.
- Art/design education appeared strongly in Gemini and also in GPT's fourth-ranked Visual-Literacy Coach, but not as Claude's or Codex's primary direction.
- Agent/platform logic was strongest in Claude's non-obvious opportunity and appeared as infrastructure in Codex and Gemini, but not as the primary GPT thesis.

4. What Stayed Stable Across Models

4.1 The apparent front door was not treated as the strongest asset

All four models moved away from treating generic single-image recognition or long-form explanation as the strongest commercial core. The language differed, but the direction was consistent.

GPT-5.6 Sol: General Visual Search’s long explanation was commercially overrated; single-image intelligence remained substitutable.

Claude Sonnet 5: General Visual Search / “Ask about anything” was the most substitutable capability and a poor flagship.

Gemini 3.6 Flash: The retained bundle combined visual analysis with web research and Journal rather than preserving generic recognition alone; the leading opportunity specialized visual guidance rather than generic identification.

Codex (GPT-5.6 Sol Medium): The value was explicitly placed in repeatable photo workflows rather than “being able to talk about a photo,” while shopping/search accuracy was judged too weak to anchor monetization.

Interpretation of the comparison: under this evidence base, rejection of the apparent core was relatively model-robust.

4.2 Reliability problems changed what the models were willing to monetize

No model ignored the documented recognition and trust failures. They handled them differently, but each either downgraded accuracy-dependent uses, limited them to lower-risk settings, or specified correction/sourcing controls as a prerequisite.

The models therefore did not simply accept the product’s current content and recognition layer at face value. They generally treated those weaknesses as constraints on where the product could credibly create value.

4.3 Recurring use became a central commercial criterion

All four answers searched for a mechanism that would generate the next use without depending on spontaneous everyday curiosity. The recurring trigger differed by model: accumulated research, personal journaling, repeated learning, campaigns, field checks, inventory intake, or client sessions. But the need for a structural repeat-use mechanism was common.

5. Where the Models Diverged

The most consequential divergence appeared at the asset-identification layer. Once the models selected different hidden assets, the rest of the strategy tree followed.

Model	Hidden asset	Recurring-value logic	Commercial consequence
GPT-5.6 Sol	Accumulated visual context across Snaps	Research corpus becomes more useful as captures accumulate	Professional visual research workspace

Claude Sonnet 5	Persistent personal memory + personalization	Archive and personal context become more valuable over time	Consumer life-journal subscription
Gemini 3.6 Flash	Visual noticing + vocabulary guidance	New objects/works create repeated learning moments	Art/design education product
Codex (GPT-5.6 Sol Medium)	Configurable photo-to-output workflow	Work itself triggers the next capture	Operational and campaign infrastructure

This is more than “different ideas.” The models disagreed about what the commercially meaningful object actually was. That disagreement appeared before buyer selection and before the final opportunity ranking.

5.1 Buyer and business-model sensitivity

The same product evidence supported four substantially different payer stories: enterprise design/retail teams, individual consumers, education institutions or learners, and brands/operational businesses. The evidence therefore did not force a stable B2C/B2B conclusion once the models moved beyond rejecting the current front door.

5.2 The “underrated feature” test exposed different attention patterns

The underrated-feature question produced especially clean divergence:

- GPT-5.6 Sol: Journal / Snap history.
- Claude Sonnet 5: AI Memory Import / personalization fields.
- Gemini 3.6 Flash: Photo to dots / poster / graphic-card generation.
- Codex (GPT-5.6 Sol Medium): Snap history and Insights.

The models were looking at the same inventory but assigned latent commercial importance to different, relatively secondary features. This is direct evidence that model choice can affect which hidden product asset receives strategic attention.

5.3 The “non-obvious opportunity” test diverged even more

- GPT-5.6 Sol: professional visual research system.
- Claude Sonnet 5: expert / agent discovery platform.
- Gemini 3.6 Flash: visual observation and vocabulary education.
- Codex (GPT-5.6 Sol Medium): field operations “capture → judgment → history → weekly evidence” infrastructure.

This question was intentionally designed to move beyond current positioning. It produced four different answers, suggesting that once the task requires strategic extrapolation rather than evaluation of an existing surface, model sensitivity rises sharply.

6. Where the Decision Chain Split

A useful way to read the outputs is as a sequence of strategic judgments:

Evidence → Asset identification → Recurring-value mechanism → User / payer → Business model → Keep / kill decisions

The strongest convergence occurred near the left side of the chain. The models largely agreed that the current front-door recognition/explanation experience was weakly differentiated and constrained by reliability issues.

The strongest divergence occurred as the models moved rightward. They selected different latent assets, and those selections generated different recurring-use mechanisms and payers. The final opportunity lists were therefore downstream expressions of earlier model-sensitive judgments.

7. Implications for AI-Assisted Product Strategy

This pilot supports a narrower conclusion than “AI strategy is unreliable” and a more useful conclusion than “different models give different answers.”

Under a fixed evidence base and fixed strategic prompt, some judgments were model-robust while positive strategy selection remained materially model-sensitive.

- 1. Agreement is not uniform across the decision.** Models may converge on diagnosis or rejection while diverging on what should replace the rejected strategy.
- 2. The main source of variance may appear before the recommendation.** In this pilot, hidden-asset identification varied. Comparing only final top-five lists would miss that earlier split.
- 3. A single-model strategy answer can conceal alternative, evidence-consistent readings.** The four outputs were not random variants of one proposal. They represented different theories of where the product’s latent commercial value resided.
- 4. Multi-model agreement can be more informative when interpreted by judgment layer.** Consensus on “what not to do” and disagreement on “what to do instead” should not be collapsed into one overall agreement score.
- 5. Model comparison does not select the final strategy.** This experiment reveals stability and variance in AI judgment. It does not establish which commercial path is correct, nor does it substitute for market validation or human strategic choice.

8. Limitations

- P01–P03 were role-play personas, not recruited users. Their reactions are exploratory evidence, not market validation.
- The evidence packet was intentionally compressed. Summary construction can introduce selection effects even when strategic interpretation is excluded.
- Codex was not run under exactly the same information pathway as the other three models: it had directly operated the product before producing its existing strategy answer. It is therefore a useful fourth comparison, but not a perfectly symmetric controlled replicate.
- Model versions, system behavior, and output tendencies can change over time. This report describes the observed runs on 18 August 2026, not permanent model characteristics.

- The pilot tests strategy generation from a fixed evidence base. It does not test real-world execution, customer willingness to pay, retention, or commercial outcomes.

9. Conclusion

Pilot #2 produced a clear asymmetry. GPT, Claude, Gemini, and Codex converged on the weakness of Chance's apparent front-door value, especially generic single-image recognition/explanation as a defensible commercial core. They did not converge on the hidden asset that should replace it.

GPT saw an accumulated visual research workflow. Claude saw persistent personal memory. Gemini saw guided visual observation and learning. Codex saw configurable, repeatable photo-to-output workflows embedded in business activity.

The practical finding is not that one model was right and the others were wrong. It is that model choice changed the strategic theory of the product while the evidence and question remained substantially fixed.

The models largely agreed on what Chance should not be, but diverged on what its hidden commercial asset actually was.

Appendix A — Neutral Evidence Packet

The following is the frozen common evidence input supplied to GPT, Claude, and Gemini. It intentionally excludes prior strategy recommendations and opportunity rankings.

Chance Pilot #2

Neutral Evidence Packet

Neutral product evidence for controlled strategy evaluation • Version 1.1 • 18 August 2026

Purpose: provide a compressed record of observed product behavior while minimizing strategic framing. This packet contains product observations, role-played persona reactions, reliability evidence, and known unknowns. It intentionally excludes prior strategy recommendations, opportunity rankings, preferred buyers, and business-model conclusions.

Evidence sources	Three structured first-use persona sessions (P01-P03) plus a persona-free capability audit.
Test environment	Primarily iPhone Mirroring. Camera-dependent actions and some text input required direct physical iPhone interaction.
External validation	Used only where explicitly noted, most importantly to verify the artwork misidentification in P02.
Excluded	Codex strategy recommendations; opportunity rankings; claims such as "strongest capability," preferred buyer, or recommended business model.

How to read this packet

“Observed facts” describe interface behavior, outputs, tested workflows, or externally verified errors. “Persona response / interpretation” describes the reaction of a specified role-play persona, not a recruited user. Persona statements are exploratory evidence and should not be treated as market validation.

1. Product Structure and Intended First-Use Experience

OBSERVED FACTS

- The observed interface was camera-centered. A primary entry point was visual capture / “Visual Search,” with additional use-specific modes or “experts.”
- Observed onboarding explained three ideas: photograph something to receive insight; swipe across different experts or perspectives; and view examples from other users.
- Across sessions, visible modes or examples included OOTD, skin analysis, personal color, nutrition, tarot, image transformation, visual search, story / entertainment modes, and social or sharing features.
- One onboarding flow asked where the user wanted to use Chance in everyday life and allowed interest selection. “Ask about anything” was preselected; “Look my best” could be added.
- A referral-code screen stated that using a friend’s code would give both users 20 credits. In one session, the account still showed 20 credits after referral was skipped.
- Japanese could be enabled through the iOS app language setting. Major interface labels then switched to Japanese, while some English strings remained.
- Camera access was required for the observed capture workflow. Notification and location prompts appeared during analyses; declining them did not prevent the observed results from being viewed.

2. P01 — Everyday Curiosity / General Consumer

Persona: Japanese university student in their early 20s, familiar with Instagram, TikTok, Google, Google Lens, ChatGPT, and Gemini. Mild curiosity about everyday objects, but unlikely to retain apps without recurring practical value.

OBSERVED FACTS

- P01 chose a free sample image containing several consumer items. Chance focused on the smoothie and identified it as Erewhon’s Hailey Bieber collaboration, “Strawberry Glaze Skin Smoothie.”
- The result was a multi-paragraph article rather than a short identification. It discussed ingredients, celebrity association, social-media popularity, and beauty positioning.
- P01 asked whether the smoothie really had beauty benefits. Chance gave a qualified answer discussing collagen peptides, hyaluronic acid, sea moss, price (about US\$20 / R\$110), lack of immediate transformative effects, and broader dietary context.
- A notification prompt appeared during analysis and a location prompt appeared during the follow-up answer. Both were declined.
- P01 stopped reading when another section on beauty / health ingredients appeared and did not explore the feed, other experts, OOTD, or tarot.
- The star icon with “20” was visible but its meaning was not explained in the tested flow.

PERSONA RESPONSE / INTERPRETATION

- Initial mental model: “AI that explains what is in a photo,” similar to Google Lens. After onboarding, the persona understood Chance as offering multiple perspectives and cultural / contextual interpretation beyond identification.
- The specific product identification and surrounding context were interesting, but the output felt too long for a light everyday question.
- The sample topic felt remote from the persona’s daily life. Notification and location requests lowered trust slightly because their relevance was unclear.
- For simple identification, the persona expected to keep using Google Lens; for detailed questioning, ChatGPT or Gemini.
- No realistic reason to reopen within seven days was identified. Expected one-month retention was low. Final characterization: “interesting but occasional utility.”

3. P02 — Art / Design Learning

Persona: Japanese university student in their early 20s beginning to study art / design, interested in learning how to look at artworks and articulate observations.

OBSERVED FACTS

- P02 used a free sample of Van Gogh's *The Starry Night* and requested more information. Chance produced a long explanation covering 1889, Saint-Rémy, memory / imagination, and cypress symbolism.
- A follow-up question about the effect of the swirling brushwork displayed "web search" and "9 sources" activity. The answer discussed impasto, gaze movement, and a pulsating appearance.
- P02 then uploaded a black-and-white woodcut. Chance identified it as Maria Bonomi, *As Meninas*, around 1964, woodcut, and generated extensive art-historical interpretation based on that identification.
- The actual work was Antonio Henrique Amaral, *O alegre anjo da loucura*, 1963, woodcut, approximately 53.5 × 80 cm. This was externally verified against gallery and marketplace records.
- Chance therefore misidentified the artist and title while estimating the medium and period closely enough to make the answer appear plausible.
- The first external verification attempt also failed because Chance's incorrect terms matched real historical material: Maria Bonomi did create 1964 woodcuts associated with a text titled "*As meninas*." The error became clear only after independent image / artist comparison.
- The work identification included an "estimated" qualifier, but the subsequent response remained long, specific, and confident.

PERSONA RESPONSE / INTERPRETATION

- Before the error was found, the persona valued Chance for converting vague visual impressions into vocabulary and points of attention, such as brushwork, carving marks, negative space, and movement.
- After the error, direct formal description retained some value, but artist identity, title, historical background, symbolism, intention, and cross-work relationships were not trusted without visible sourcing.
- The persona preferred a short first layer showing likely identification, what to notice, and why it may be interesting, with longer detail available on demand.
- Expected future workflow after the error: use Google Lens or official museum information to establish identity first, then potentially use Chance to expand viewing questions or perspectives.
- Use would likely be selective rather than frequent, such as one interesting work among several museum objects.

4. P03 — Fashion / Vintage Clothing

Persona: Japanese university student in their early 20s who regularly browses vintage, flea-market, antique, secondhand, and independent fashion stores.

OBSERVED FACTS

- P03 selected "OOTD fit check." The mode asked for an outfit image and a use context. "Campus" was selected.
- The image showed a black-and-white striped dress and black closed shoes. Chance initially suggested trying a more suitable photo, then produced a result titled "Monotone Summer."
- The output included style tags, general outfit commentary, praise for mixed stripe widths, and suggestions such as silver accessories or a hat.
- Chance recommended sandals and approximately 23°C weather. The photographed footwear appeared to be black closed shoes, not sandals.
- The selected "Campus" context was not clearly reflected in the final recommendation.
- A notification prompt appeared during analysis and was declined.

PERSONA RESPONSE / INTERPRETATION

- The persona understood the distinction from Google Lens as interpretation / styling rather than identification.
- The workflow was easy to understand, but the styling advice felt generic. Shoe misclassification and weak use of the selected context reduced confidence in personalization.
- The persona was more interested in help noticing era, material, construction, unusual details, historic references, or whether one specific accessory works than in receiving a single "correct" outfit.
- Expected use was occasional. If repeated outputs remained generic, the persona expected to prefer showing an image to ChatGPT or another conversational AI.

5. Persona-Free Capability Audit

After the persona sessions, testing switched to a broader capability audit. The audit inspected accessible modes, outputs, creation tools, history / memory surfaces, sharing flows, and related product behavior beyond a normal first-use path.

Capability	Status	Input	Observed / displayed output
General visual search	Tested	Photo + intent such as “tell me” / “where can I buy?”	Explanation, attributes, sources, product candidates, prices, poster-like visual
Plant recognition	Tested	Plant photo	Species estimate, explanation, care follow-up, 19 sources
OOTD	Tested	Outfit photo + context	Outfit title, tags, commentary, improvement suggestions
Photo to dots	Tested	Photo	Decorative card retaining photo, extracted palette, editable layout
Skin Analysis	Screen only	Face selfie	Example showed skin-state visualization + routine suggestions
Glow-up	Screen only	Selfie	Example showed before / after image changes to hair, makeup, lighting
Personal Color	Screen only	Selfie	Palette + style card; shown as referral-unlocked
Nutrition Lens	Screen only	Food photo	Example calorie / protein / carbohydrate overlay
Today's Tarot	Screen only	Question + surrounding-object photo	Photo-integrated tarot card + interpretation; shown as once daily
Minion Buddy	Screen only	Person / animal photo	Example composited Minions into photo
Day Flow	Screen only	Current-situation photo	BeReal-like public post with time / activity
Snap Shop	Screen / locked	Product photo	Described as price, retailer, similar-product search
Agent discovery	Tested	Search / filter / save	20+ analysis, divination, and transformation modes
Agent creation	Screen only	Discord participation required	No in-app creation screen found; Discord redirect
Snap history / Journal Insights	Tested	Execution of modes	Automatic card history of analyses / transformations
	Tested	Multiple Snaps	Weekly calendar / recap, Patterns, Taste Profile, nutrition tracker
Weekly Recap	Tested	Dates / selected Snaps	Static collage using about three images
AI Memory Import	Tested	Pasted self-summary from another AI	Adds text to Chance memory; copy-paste, not observed API sync
Personalization	Tested	Tone, occupation, interests, values, custom instructions	Adjusted framing / wording of analyses
Sharing	Tested	Result card	Full-page share, link copy, image save, Instagram, WhatsApp, etc.
Chance feed post	Screen only	Result + Post Snap	Publishes result to Chance feed; not executed
Text-to-speech	Screen only	Result text	Headphone control for read-aloud

6. Journal, Memory, and Longitudinal Product Behavior

OBSERVED FACTS

- Completed analyses and transformations were automatically saved as “Snaps” in the Journal. Both free samples and user-generated analyses were retained.
- In one session, Chance generated a cross-image statement from two saved images and showed an “one more image” progress indicator encouraging a third interaction.
- Insights included a weekly calendar, weekly usage display, “Patterns,” “Taste Profile,” nutrition tracking, and a daily tracker. At the inspected point, “discovered patterns” showed zero.
- Weekly Recap accepted selected dates / Snaps and produced a static collage using approximately three images.
- AI Memory Import allowed a user to paste a self-summary generated by another AI into Chance memory. The observed flow was manual copy-and-paste rather than direct API connection.
- Personalization fields included tone, occupation, interests, values, and custom instructions. These could affect how analysis text was framed.

7. Generation, Sharing, and Social Flows

OBSERVED FACTS

- Photo to dots was executed. It produced a decorative card that retained the original image, extracted a color palette, and offered layout editing. The observed output was not literal dot / pixel conversion.
- General visual search could also generate a poster-like visual based on the analyzed subject, including subject naming and color / layout treatment.
- Displayed but unexecuted transformation examples included cyberpunk wallpaper, artwork collage, Glow-up, Minion compositing, tarot-card generation, and story / horror outputs.

- Result cards could be shared as a full page, copied link, saved image, or through destinations including Instagram and WhatsApp.
- A “Post Snap” flow could publish a result into Chance’s own feed; posting was not executed.
- Day Flow appeared to support public posts combining a current image, time, and activity context. The audit did not establish feed activity, audience scale, or network effects.

8. Reliability and Trust Evidence

Area	Observed issue	Concrete evidence
Artwork attribution	Incorrect identification	Artist and title were wrong; long downstream interpretation was generated from the incorrect attribution.
Fashion OOTD	Object / context mismatch	Black closed shoes were described as sandals; “Campus” selection had little visible effect.
Shopping / product match	Large price / match mismatch	Known R\$280 cushion was matched with approximately R\$19.90–42.00 items, apparently mixing related products / inserts with an exact match.
Object context	Unverified narrative	Cushion was linked to the Brazilian cactus “Coroa de Frade” and artisan-production claims not established during the test.
Plant target selection	Target-selection issue	Automatic selection included background, container, and support; no post-submission correction control was found.
Sources	Uneven / unclear attribution	Source lists included uneven-quality domains such as Facebook; sources were not consistently seen attached claim-by-claim.

9. Repeated First-Use Patterns Across P01–P03

ROLE-PLAY RESPONSE PATTERNS

- Across the three role-play sessions, each persona articulated a basic mental model of the product and identified at least one useful or interesting aspect.
- P01 identified no realistic reason to reopen within seven days; P02 described future use as selective after the attribution error; P03 described expected use as occasional unless outputs became more specific.
- Notification permission was declined in all three persona sessions in which it appeared.
- In their stated post-session preferences, the personas continued to name Google Lens, ChatGPT, and / or Gemini for adjacent tasks.
- P01 and P02 each described the output length as heavy for repeated in-the-moment use.
- P02 reported lower trust after the artwork attribution error. P03 reported lower confidence in personalization after the footwear mismatch and weak reflection of the selected “Campus” context.
- These are structured role-play reactions, not observed behavior from recruited users.

10. Known Unknowns and Untested Areas

- Full Snap Shop behavior and exact-match shopping accuracy.
- Output quality and accuracy for Skin Analysis, Glow-up, Personal Color, Nutrition Lens, and most entertainment agents.
- Quality consistency across the full agent catalog and the external / Discord-based creator workflow.
- Real activity, network effects, or retention value of the Chance feed and Day Flow.
- Long-term usefulness of Patterns, Taste Profile, Insights, imported memory, and longitudinal personalization.
- Exact credit economics, subscription structure, paid conversion, and willingness to pay.
- Whether observed recognition errors are systematic, category-specific, or occasional.
- Whether source retrieval, ranking, and claim-level attribution are consistent across modes.

- Whether location materially improves outputs and whether notification prompts produce useful re-engagement.
- Real-user retention, repeat frequency, and behavior outside the three simulated personas.

11. Neutral Capability Inventory

The audit recorded the following observed or displayed capability types within the product:

- Visual recognition and explanation; object / plant analysis; fashion analysis; shopping lookup; conversational follow-up; web-assisted research.
- Image transformation and poster / card generation; beauty simulation; nutrition analysis; themed entertainment / divination outputs.
- Journal / Snap history; cross-Snap summaries; weekly recap; Pattern / Taste Profile surfaces; imported AI memory; user personalization.
- Sharing to external destinations; internal social posting; agent discovery; read-aloud.

This inventory describes observed product structure only. It does not identify a core asset, preferred market, buyer, business model, or strategic priority.

Appendix B — Common Strategy Prompt

Use only the attached Evidence Packet as your evidence base.

Do not browse the web, search for additional information, or use outside knowledge about Chance. Do not assume capabilities, customer behavior, monetization, or product characteristics that are not supported by the Evidence Packet.

The persona responses in the Evidence Packet are role-played exploratory evidence, not findings from recruited user research. Treat them accordingly.

Now switch from capability discovery to business strategy.

Based primarily on the observed evidence in the packet, identify the strongest commercial opportunities for Chance.

Do not preserve the current consumer positioning by default.

You may discard most current features.

Do not assume text explanation or visual search is the core asset.

Prefer opportunities that make distinctive use of capabilities actually observed in the product, especially where multiple capabilities combine into a workflow.

For each opportunity, explain:

- what capability or combination of capabilities creates the value;
- who uses it;
- who pays;
- why the use is recurring;
- why generic multimodal AI alone is not sufficient;
- what current Chance features would be kept;
- what would be discarded;
- what would need to change in the product;
- the strongest reason the idea might fail.

Rank the top 5 opportunities.

Finally answer:

- If you could keep only one capability or capability bundle from the current product, what would you keep?
- Which current feature is commercially underrated?
- Which current feature is commercially overrated?
- What is the strongest opportunity that is not obvious from Chance's current positioning?

For each major strategic judgment, identify the specific observations from the Evidence Packet that materially influenced the judgment. Clearly distinguish direct evidence from your inference.

Base the analysis on the Evidence Packet, not on generic AI business patterns.